When Retrieval Quality Decouples from Faithfulness: A Pre-Registered Audit of RAG Evaluation

Anonymous Authors
Submission under double-blind review

Abstract

Retrieval-augmented generation (RAG) papers often treat better retrieval quality as evidence that answers are more faithful. We show that the faithfulness scorer can be the under-identified variable: on 7,500 fixed generations, one intervention has effect sizes 0.011, 0.032, and 0.140 under DeBERTa, a second NLI scorer, and a local RAGAS-style judge. A 99-example two-rater human calibration reaches Cohen's $\kappa = 0.774$, but alignment with adjudicated support is weak for DeBERTa and moderate for the other scorers (Spearman $\rho = 0.140, 0.380, 0.441$); bootstrap intervals overlap, so the ordering is suggestive. We then use a pre-registered refinement-paradox stress test to audit identifiability once retrieval similarity, context-set structure, evaluator choice, human calibration, threshold transfer, and cost are separated. At $n = 500$ across five seeds, the DeBERTa-measured SQuAD/Mistral pilot effect shrinks to 1.1 faithfulness points, with only 1/5 seeds significant. A matched HIGH/LOW context-coherence test does not support a causal CCS interpretation (HIGH–LOW $= -0.002$, $p = 0.628$, 95% CI $[-0.022, 0.017]$); HIGH-CCS contexts hallucinate more often (16.5% vs. 9.0%). Threshold transfer is uneven, and HCPC-v2, CRAG, and RAPTOR-2L yield no single deployment ordering once tail latency and indexing cost are reported; Self-RAG is documented separately as a harness-mismatched appendix baseline. Coherent uninformative noise hurts less than random off-topic noise. We propose a minimum reporting standard for RAG faithfulness claims (paired query-level contrasts, retrieval/context statistics, multiple scorers, human calibration on disagreements, threshold-transfer evidence, hallucination rates, p99 latency, indexing cost) and release ControlledRAG as the audit protocol that operationalizes it.

1 Introduction

RAG evaluation has a familiar shortcut: a retrieval score improves, one automatic faithfulness metric moves in the desired direction, and the paper infers that answer support has improved. Our audit finds that the scorer can be the fragile link. On the same 7,500 generated answers, the same intervention has effect sizes 0.011, 0.032, and 0.140 under DeBERTa, a second NLI scorer, and a local RAGAS-style judge. In a 99-example two-rater calibration ($\kappa = 0.774$), DeBERTa has weak rank alignment with adjudicated support labels, while the second NLI scorer and RAGAS-style judge are only moderately aligned; their confidence intervals overlap, so the ordering is suggestive rather than definitive. The durable finding is not that one scorer is best, but that RAG faithfulness claims are not identified by a single observer.

This paper audits that gap. A generator receives a context set, not a list of independent passage scores, and a faithfulness conclusion depends on the observer used to score support. We use the refinement paradox as a stress test: the pilot finding that more aggressive per-passage refinement can raise retrieval similarity while lowering measured answer faithfulness. The goal is not to promote that phenomenon as large or stable. It is to ask a sharper methodological question:

> *Which RAG faithfulness claims remain identifiable once local retrieval similarity, context-set structure, evaluator choice, human calibration, threshold transfer, and cost are separated?*

The answer is sobering. The strongest intuitive story – that context coherence explains faithfulness failures – fails under a direct control. We construct HIGH- and LOW-CCS context sets matched on mean query-passage similarity; the coherence manipulation is large, but faithfulness does not improve. This failure is the point. It shows how a plausible retrieval-time quantity can look meaningful in aggregate tables and then lose its apparent effect once the confounder is isolated.

The wider audit reaches the same conclusion from other angles. The pilot DeBERTa effect becomes small and seed-sensitive at $n = 500 \times 5$. The same generations imply different intervention magnitudes under different scorers. A small two-rater human calibration clarifies that the scorer disagreement is not cosmetic: no automatic scorer has strong human alignment in this sample, although the RAGAS-style judge has the highest rank correlation. Thresholds tuned on one dataset transfer unevenly. Baseline comparisons change when p99 latency and indexing cost are included.

Recommended minimum reporting standard for RAG faithfulness claims.

1. Paired query-level contrasts with bootstrap intervals.
2. Local query-passage similarity and at least one set-level context statistic.
3. Multiple faithfulness scorers.
4. A small human calibration whenever automated metrics disagree.
5. Threshold transfer across datasets.
6. Hallucination rates with Wilson intervals.
7. Mean latency, p99 latency, and offline indexing cost.

Contributions.

- We show that faithfulness scoring is metric-fragile on fixed generations, and calibrate the disagreement with a 99-example two-rater human audit.
- We provide a pre-registered scale audit of a pilot refinement-paradox cell, replacing a small-sample headline with $n = 500 \times 5$ seed evidence.
- We falsify the strong CCS interpretation with a matched-similarity intervention, while preserving a narrower diagnostic role via a coherent-versus-random noise test.
- We audit threshold transfer and cost-aware baselines, and document the Self-RAG artifact separately as a harness-mismatched appendix baseline.
- We package these checks as ControlledRAG and propose them as the minimum reporting standard for RAG faithfulness claims; anonymized code and data ship with the supplement.

2 Background and Problem Setup

RAG faithfulness. RAG systems retrieve passages and condition a generator on the resulting context [Lewis et al., 2020, Karpukhin et al., 2020]. A generated answer is faithful when its substantive claims are supported by the retrieved context, not merely when the answer is true in the world. This distinction matters because a model can answer from parametric memory even when retrieval failed, or can produce unsupported details despite apparently relevant passages [Maynez et al., 2020, Ji et al., 2023].

Local retrieval similarity. Most retrieval metrics are local: they score each query-passage pair with embedding similarity, cross-encoder relevance, or reranking rank. Local scores are useful, but they do not identify answer support. Three individually relevant passages may be redundant, contradictory, or missing a bridge fact. One lower-scoring passage may contain the answer.

Context coherence and CCS. The Context Coherence Score (CCS) is a retrieval-time statistic over a retrieved context set C: $\mathrm{CCS}(C) = \mathrm{mean}_{i<j} \cos(e_i, e_j) - \mathrm{sd}_{i<j} \cos(e_i, e_j)$, where e_i and e_j are passage embeddings. The mean term rewards mutual topical compatibility, while the standard-deviation penalty reduces scores for uneven sets with one coherent sub-cluster and one outlying passage. CCS asks a set-level question that local relevance does not: do the passages look mutually compatible? The audit treats CCS as a diagnostic candidate. It does not assume that CCS causes faithfulness, and the matched test in Section 4 rejects that strong interpretation.

Identifiability. We use "identified" in a practical evaluation sense. A RAG claim is not identified if another unreported factor could plausibly explain the reported effect. A local retrieval improvement does not identify support if context-set structure is uncontrolled. A single automatic metric does not identify an effect size if other faithfulness observers disagree. A tuned threshold does not identify a reusable policy if off-diagonal transfer is poor. A method does not identify deployment value if p99 latency and indexing cost are omitted.

Why the usual evidence is risky. The risky combination is small samples, one scorer, threshold tuning on the reporting domain, no human calibration, and partial cost reporting. Each choice can be reasonable in isolation. Together, they let a paper report a clean faithfulness story without separating the variables needed to support it.

3 Pre-Registered Audit Protocol: ControlledRAG

The experiments are a controlled falsification suite. Each test isolates a claim that can look convincing in a conventional aggregate RAG table, then asks whether the claim survives once the relevant confounder is separated.

Pre-registration and statistics. Sample sizes, seeds, primary statistics, and decision rules were frozen before execution. Paired continuous contrasts use Wilcoxon signed-rank tests and Cohen's d_z. Means use 10,000-resample bootstrap confidence intervals. Hallucination rates use Wilson intervals. Thresholds are frozen except in the threshold-transfer experiment, where tuning and evaluation datasets are explicitly separated. No paid API calls are used in the revision results.

Datasets, generators, and scorers. The core revision evidence uses short-answer extractive QA, primarily SQuAD, PubMedQA, HotpotQA, NaturalQuestions, and TriviaQA [Rajpurkar et al., 2016, Jin et al., 2019, Yang et al., 2018, Kwiatkowski et al., 2019, Joshi et al., 2017]. The main scaled cell is SQuAD with Mistral-7B-Instruct [Jiang et al., 2023]. Faithfulness is scored with DeBERTa-v3 NLI [He et al., 2021], a second NLI model, and a local RAGAS-style Mistral-as-judge rubric [Es et al., 2024]. The local judge is useful for triangulation but may share family-level biases with the generator; we treat that as a limitation rather than a source of authority.

Human calibration design. To ground the scorer disagreement, two raters annotate 99 examples sampled from the fixed-generation multi-metric cell. Labels are `supported`, `partially_supported`, and `unsupported`; adjudicated labels are mapped to $1, 0.5, 0$ for rank correlations. The instructions require raters to judge support using only the provided passages.

Table 1: ControlledRAG audit controls. Each control can be reused with other datasets, retrievers, generators, and faithfulness metrics.

Control	Claim it tests	Failure mode exposed	Primary evidence
Similarity-matched intervention	CCS improves faithfulness after local relevance is fixed	A set statistic can lose its apparent effect once mean query-passage similarity is controlled	Table 2
Multi-seed scale correction	A pilot effect remains stable under repeated sampling	Small effects can be seed-sensitive even when the pilot direction is plausible	Table 3
Fixed-generation multi-metric and human-calibrated scoring	One automatic judge identifies the intervention effect size	The same outputs can imply different magnitudes, and human alignment can be only moderate	Tables 4–6
Threshold-transfer matrix	A tuned threshold is a reusable policy	In-domain tuning can hide weak off-diagonal transfer	Table 7
Coherent-vs-random noise separation	Coherence is just generic retrieval quality	Same-topic uninformative context and random context can have different slopes	Table 8
Cost-aware baseline audit	A method is deployment-preferable	Dataset wins can disappear once hallucination, p99 latency, and indexing cost are included	Fig. 2; Table 9

ControlledRAG as a minimum reporting standard. ControlledRAG is prescriptive: the seven-item list in Section 1 is the minimum admissible evidence for a RAG faithfulness claim. A claim drawn from one retrieval score, one scorer, one tuned threshold, one in-domain dataset, or one favorable baseline row is flagged as under-identified. Reusing ControlledRAG reduces under-identification; it does not certify faithful outputs or crown a single scorer.

4 Matched-Similarity Causal Test

The strong CCS hypothesis predicts that, at fixed local retrieval similarity, higher context coherence should improve faithfulness. We test that directly. For each SQuAD query, the protocol retrieves candidate passages and chooses a HIGH-CCS and LOW-CCS three-passage set whose mean query-passage similarity differs by at most 0.02. The construction succeeds: mean absolute similarity gap is 0.006 and mean CCS gap is 0.533 across 200 matched pairs.

Table 2: Matched-similarity intervention. Higher CCS does not improve faithfulness at fixed mean query-passage similarity.

Setting	n	HIGH faith	LOW faith	HIGH–LOW	p	95% CI
SQuAD / Mistral / DeBERTa	200	0.636	0.639	-0.002	0.628	[-0.022, 0.017]

The binary hallucination result points in the wrong direction for the strong claim: HIGH-CCS contexts hallucinate more often than LOW-CCS contexts (16.5% vs. 9.0%). We therefore reject the causal CCS interpretation tested here. CCS may still be diagnostically useful, but it is not sufficient to identify answer support once local retrieval similarity is held fixed.

A span-presence check suggests a simple reason for the wrong-sign result. We normalize the gold answer and ask whether it appears anywhere in the three retrieved passages. LOW-CCS contexts contain the answer more often than HIGH-CCS contexts (27.0% vs. 17.5%; exact McNemar $p = 0.011$); the four-way count is 19 pairs with the answer in both, 35 in LOW only, 16 in HIGH only, and 130 in neither.

Thus the HIGH-CCS sets do not merely provide an equally answer-bearing but more coherent scaffold; they more often omit the answer span. A plausible reading is that coherent-looking context can still be incomplete, encouraging fluent unsupported answers when the needed span is absent [Maynez et al., 2020].

We report this as a hypothesis suggested by the data, not as a tested mechanism. The null does not imply context structure is irrelevant; it shows coherence alone is insufficient, and ControlledRAG therefore audits answer-bearing evidence, evaluator choice, transfer, and cost as separable axes (§3).

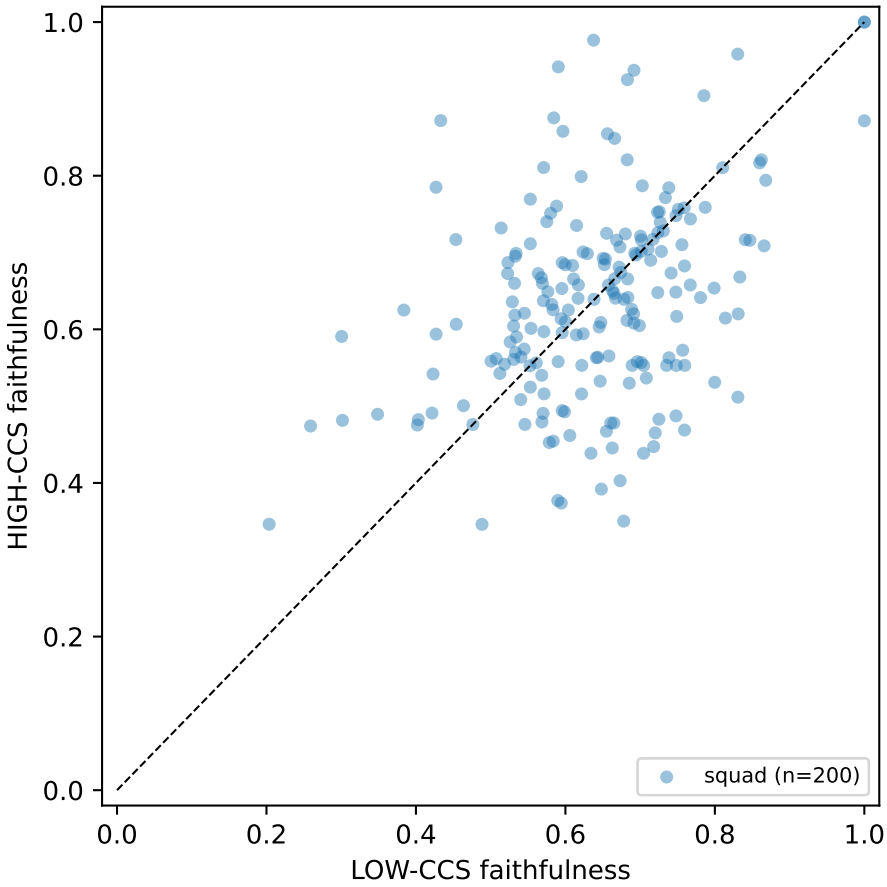

Figure 1: Paired HIGH–LOW CCS faithfulness differences under the matched-similarity intervention. The diagonal line represents equal HIGH and LOW faithfulness; points above the line are pairs where HIGH-CCS helped, and points below it are pairs where it hurt. The distribution straddles the diagonal, so the controlled test rejects the strong CCS hypothesis rather than supporting it.

5 Scaling the Pilot Paradox

The pilot refinement paradox is interesting only if it survives adequate sampling. We therefore rerun the SQuAD/Mistral cell at $n = 500$ over five seeds and three conditions: dense baseline, HCPC-v1 refinement, and HCPC-v2 gated refinement. Under the DeBERTa scorer, the baseline-minus-HCPC-v1 effect is 1.1 faithfulness points, and only one of five seeds reaches $p < 0.05$ for the paired contrast.

Table 3: Scaled SQuAD/Mistral headline cell. The pilot direction survives as a small DeBERTa-measured non-monotonicity, not as a large stable effect.

Condition	n	Faithfulness	Hallucination	Retrieval similarity
Baseline	2500	0.661	0.146	0.532
HCPC-v1	2500	0.650	0.147	0.569
HCPC-v2	2500	0.661	0.143	0.532

The result should be read as a scale correction. HCPC-v1 raises retrieval similarity while lowering DeBERTa faithfulness, but the magnitude is small and seed-sensitive. HCPC-v2 recovers the HCPC-v1 drop on average, yet this does not establish a general solution to hallucination or a global method ranking.

6 Faithfulness Scoring Is Metric-Fragile

The same generated answers can tell different stories depending on the faithfulness observer. We therefore rescore the full 7,500-row scaled cell with DeBERTa-v3, a second NLI model, and a local RAGAS-style judge while holding the generations fixed.

For correlation uncertainty, we use 10,000 paired-index bootstrap resamples with a fixed seed and percentile 95% intervals. Each resample preserves the pairing between the two scores being correlated; human-label correlations resample the adjudicated ordinal labels and automatic score together.

Table 4: Identical generations produce different intervention magnitudes under different faithfulness judges.

Condition	n	DeBERTa	Second NLI	RAGAS-style judge
Baseline	2500	0.661	0.350	0.730
HCPC-v1	2500	0.650	0.318	0.590
HCPC-v2	2500	0.661	0.351	0.728
Baseline–v1	–	0.011	0.032	0.140

Table 5: Metric agreement on the same 7,500 generations. Brackets give paired bootstrap 95% CIs.

Metric pair	Pearson r [95% CI]	Spearman ρ [95% CI]
DeBERTa / second NLI	0.259 [0.240, 0.278]	0.265 [0.245, 0.285]
DeBERTa / RAGAS-style judge	0.182 [0.162, 0.202]	0.212 [0.191, 0.233]
Second NLI / RAGAS-style judge	0.674 [0.661, 0.686]	0.651 [0.638, 0.665]

This is the central positive result. The DeBERTa-measured effect is small; the second NLI effect is moderate; the RAGAS-style judge effect is roughly an order of magnitude larger. Because DeBERTa and the RAGAS-style judge correlate weakly on identical generations, the disagreement is not a monotone rescaling. The §6.1 human calibration shows none of the three scorers is strongly aligned with adjudicated support, so evaluator choice changes the empirical conclusion and no single automatic observer in this audit cleanly adjudicates the choice.

6.1 Human Calibration of Metric Disagreement

To calibrate the scorer disagreement, two raters annotate 99 examples from the fixed-generation cell. The raw agreement is 0.919 and Cohen's κ is 0.774. Adjudicated labels are 75 supported, 20 partially supported, and 4 unsupported. Spearman correlation with the adjudicated ordinal labels is weak for DeBERTa, moderate for the second NLI scorer, and moderate for the RAGAS-style judge.

Table 6: Human calibration audit ($n = 99$, two raters, adjudicated labels). Brackets give paired bootstrap 95% CIs.

Metric	Spearman ρ [95% CI]	Pearson r [95% CI]	Notes
DeBERTa-v3 NLI	0.140 [-0.047, 0.314]	0.145 [-0.016, 0.296]	Weak alignment
Second NLI	0.380 [0.208, 0.534]	0.385 [0.225, 0.523]	Moderate alignment
RAGAS-style judge	0.441 [0.249, 0.624]	0.338 [0.121, 0.542]	Moderate alignment

The human calibration does not settle faithfulness evaluation generally. It does rule out a comfortable interpretation in which any one automatic metric can be treated as obviously correct. In this sample, the RAGAS-style judge is closest to human rank order, but only moderately; DeBERTa is poorly aligned. The

pairwise differences in rank alignment between scorers have overlapping bootstrap intervals, so the relative ordering should be read as suggestive rather than definitive. The safe conclusion is that metric disagreement must be reported, and small single-metric effects should not carry broad claims.

7 Threshold Transfer

Thresholded retrieval policies can look strong when the threshold is tuned and evaluated on the same dataset. We test transfer by tuning τ on each of five datasets and evaluating on all five. For evaluation dataset d and threshold τ, recovery is

$$\text{Recovery}(d, \tau) = \frac{\text{Faith}_{\text{gate}}(d, \tau) - \text{Faith}_{\text{HCPC-v1}}(d)}{\text{Faith}_{\text{baseline}}(d) - \text{Faith}_{\text{HCPC-v1}}(d)}.$$

Scores can exceed 1 when the gated policy exceeds the baseline relative to HCPC-v1, and can be negative when transfer is worse than HCPC-v1 under this normalization.

Table 7: Threshold-transfer recovery score. SQuAD, PubMedQA, and NaturalQuestions show large in-domain versus off-diagonal gaps; TriviaQA and HotpotQA do not. *Worked example.* On PubMedQA, baseline $= 0.559148$, HCPC-v1 $= 0.562581$, gated at τ=0.4 $= 0.557595$ (table rounded to three decimals); recovery $= (0.557595 - 0.562581)/(0.559148 - 0.562581) = 1.452$. Because HCPC-v1 happens to outperform the dense baseline on PubMedQA, this above-1 score means the gated policy fell 45.2% of |baseline − HCPC-v1| below the baseline; a negative score means the gated policy underperforms HCPC-v1 under this normalization.

Tune dataset	Diagonal recovery	Off-diagonal mean	Gap
PubMedQA	1.452	0.451	1.002
NaturalQuestions	1.401	0.211	1.191
SQuAD	0.823	-0.141	0.963
TriviaQA	0.464	0.698	-0.234
HotpotQA	-0.112	0.093	-0.205

The result is uneven transfer, not a universal failure. Some datasets show a large in-domain advantage, while others do not. Reporting only the tuned diagonal would make the policy look more stable than the off-diagonal evidence supports.

8 What Survives: Coherence vs. Similarity

The matched-similarity test rejects the strong CCS interpretation, but it should not be overread as evidence that context-set structure is useless. We compare two kinds of uninformative context: random off-topic noise and same-topic passages that preserve some contextual coherence while omitting the answer. If coherence were only generic retrieval quality under another name, the slopes should look similar.

Table 8: Noise slope response. Coherence-preserving uninformative noise hurts less than random off-topic noise at matched rate.

Condition	Faith slope	Similarity slope	Drop at full noise
Random noise	-0.069	-0.481	0.052
Coherent uninformative noise	-0.043	-0.113	0.042

This supports a narrower claim. Coherence carries diagnostic signal beyond mean query similarity, but it does not identify support by itself. A coherent context set can still omit the answer, and a low-coherence set can still contain the necessary evidence.

9 Cost-Aware Baseline Audit

This section is not a leaderboard. It asks whether a deployment-level preference remains identifiable once p99 latency and offline indexing cost are reported alongside faithfulness and hallucination. We compare HCPC-v2, CRAG [Yan et al., 2024], and RAPTOR-2L [Sarthi et al., 2024] in a matched short-answer harness. Self-RAG [Asai et al., 2024] is reported separately in the supplement: its released fine-tuned Llama-2-7B checkpoint, 8-bit inference path, reflective decoding behavior, and entailment-style scoring make that comparison harness-conditional.

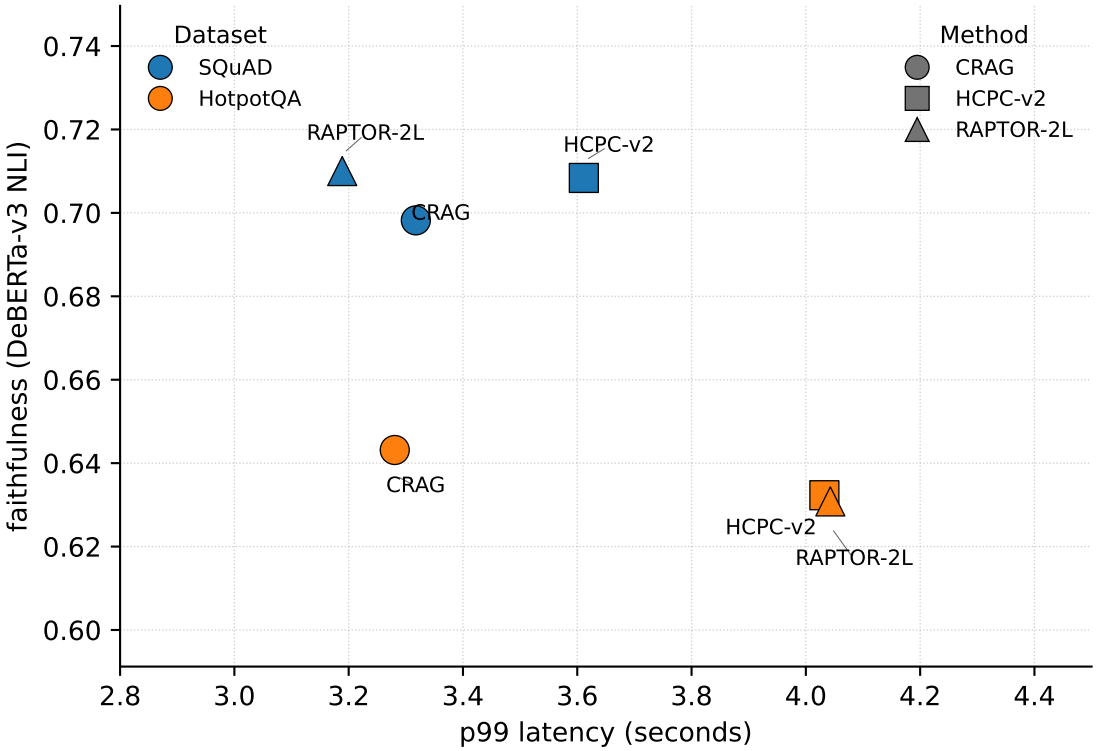

Figure 2: Cost-aware view of the matched-harness baseline audit. The x-axis is p99 latency in seconds; method is encoded by marker shape and dataset by color. No matched-harness method has a single deployment ordering across datasets once faithfulness, hallucination behavior, tail latency, and offline indexing costs are considered. The harness-mismatched Self-RAG numbers are reported separately in the supplement.

Table 9: Head-to-head baseline comparison, $n = 200$ per dataset/condition. Latency is milliseconds.

Dataset	Method	Faith	Halluc.	Mean lat.	p99 lat.
SQuAD	CRAG	0.698	0.120	1021	3318
SQuAD	HCPC-v2	0.708	0.125	1349	3612
SQuAD	RAPTOR-2L	0.710	0.125	1296	3189
HotpotQA	CRAG	0.643	0.105	1473	3281
HotpotQA	HCPC-v2	0.632	0.130	1841	4032
HotpotQA	RAPTOR-2L	0.631	0.130	1906	4043

HCPC-v2 is competitive on SQuAD, where RAPTOR-2L is slightly higher in faithfulness and lower in p99 latency. On HotpotQA, CRAG has higher faithfulness, lower hallucination, and lower latency. Self-RAG numbers are in the supplement; a fair evaluation requires a scoring protocol that accounts for reflective decoding rather than treating every extra token as ordinary answer content. The preferred matched-harness method changes once dataset and cost axes are included, so a single-row deployment claim is not identified

by this evidence.

Table 10: RAPTOR cost audit across three datasets ($n = 100$ each). Latency columns are milliseconds; indexing columns are wall-clock seconds.

Dataset	Faith	Halluc.	p50 lat.	p99 lat.	Dense idx.	RAPTOR idx.
SQuAD	0.789	0.050	1190	4943	8.333 s	161.203 s
PubMedQA	0.560	0.290	3901	6812	1.177 s	108.264 s
HotpotQA	0.617	0.210	1973	4274	7.140 s	99.740 s

10 Discussion

The collapse is informative. The audit weakens the simple refinement-paradox story: the pilot effect becomes small under DeBERTa, the matched CCS test rejects the strong coherence interpretation, and the baseline ranking changes with dataset and cost surface. That is the point. Standard RAG reporting under-identifies faithfulness claims.

What can and cannot be claimed. Local retrieval quality and answer faithfulness can decouple; CCS is diagnostic, not causal; metric choice materially changes intervention magnitude; human calibration gives only moderate alignment for the closest automatic scorer; threshold transfer is uneven; and baseline comparisons require hallucination, p99 latency, and indexing cost. The audit does not show that HCPC-v2 solves hallucination, that the DeBERTa paradox is large and stable at scale, that Self-RAG is globally weak, or that any automatic scorer is a universal faithfulness metric. Future claims that omit an axis of Section 3 should be marked under-identified.

11 Related Work

RAG combines parametric generation with retrieved evidence [Lewis et al., 2020, Karpukhin et al., 2020]; selective context compression [Xu et al., 2024] reduces noise without changing the audit question we pose about the resulting context set. Automatic faithfulness scoring builds on NLI entailment [Honovich et al., 2022, Laban et al., 2022], atomic-claim decomposition [Min et al., 2023], RAG-specific rubrics [Es et al., 2024], and LLM-as-a-judge protocols [Liu et al., 2023]; the literature observes scorer disagreement across systems, but to our knowledge this paper is the first to operationalize the disagreement on a fixed 7,500-row generation cell with paired bootstrap intervals and a two-rater human calibration in one protocol. Corrective retrieval, reflective generation, and hierarchical summary retrieval [Yan et al., 2024, Asai et al., 2024, Sarthi et al., 2024] are method instances; our contribution is not a new retriever or a universal ranking but ControlledRAG, an audit protocol so these and future methods can be evaluated on identified claims.

12 Limitations and Ethics

Scope. The main evidence concerns short-answer extractive QA on SQuAD/Mistral and should not be used to claim the same behavior in long-form synthesis, agentic workflows, or domain-specific retrieval without additional controls. The matched-similarity test rejects the strong CCS interpretation in that setting but does not exhaust all context-structure hypotheses; threshold transfer is also uneven, so thresholded policies should report off-diagonal behavior.

Human calibration and judge bias. The human evaluation has 99 examples, two raters, and adjudicated labels: it calibrates scorer disagreement in the audited cell, but does not definitively validate any faithfulness metric. The RAGAS-style judge is local and zero-dollar, but may share family-level biases with the generator; the stronger conclusion is metric fragility, not that the judge is universally correct.

Self-RAG and large-model reproduction. Self-RAG is a harness-mismatched appendix baseline (different backbone, quantization, and reflective decoding objective); its numbers should not be read as a global statement about Self-RAG. The pilot's 70B and 120B observations are not re-verified under the zero-dollar revision constraint and are therefore not treated as central evidence.

Ethics. The audit is intended to reduce overclaims about faithful RAG systems. The main risk is misuse: a single metric, a single threshold, or one favorable baseline row could be cherry-picked to market a system as grounded. We therefore emphasize per-query logs, multiple metrics, human calibration, negative results, and cost-aware reporting.

13 Conclusion

Better local retrieval, higher context coherence, a single faithfulness metric, an in-domain threshold, or one favorable baseline row does not identify RAG answer support. The central durable contribution is the fixed-generation metric-fragility finding—the same outputs imply materially different intervention magnitudes under different scorers, and human calibration shows none can be treated as unqualified ground truth. We release ControlledRAG as the minimum standard of evidence for making RAG faithfulness claims; a paper that omits paired contrasts, multiple scorers, human calibration on disagreements, threshold transfer, p99 latency, or indexing cost is, by this standard, under-identified.

References

Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. Self-RAG: Learning to retrieve, generate, and critique through self-reflection. In *International Conference on Learning Representations*, 2024.

Shahul Es, Jithin James, Luis Espinosa-Anke, and Steven Schockaert. RAGAS: Automated evaluation of retrieval augmented generation. In *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics: System Demonstrations*, pages 150–163, 2024.

Pengcheng He, Xiaodong Liu, Jianfeng Gao, and Weizhu Chen. DeBERTa: Decoding-enhanced BERT with disentangled attention. In *International Conference on Learning Representations*, 2021.

Or Honovich, Roee Aharoni, Jonathan Herzig, Hagai Taitelbaum, Doron Kukliansy, Vered Cohen, Thomas Scialom, Idan Szpektor, Avinatan Hassidim, and Yossi Matias. TRUE: Re-evaluating factual consistency evaluation. *NAACL*, 2022.

Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Yejin Bang, Andrea Madotto, and Pascale Fung. Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12):1–38, 2023.

Albert Q. Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, et al. Mistral 7B. *arXiv preprint arXiv:2310.06825*, 2023.

Qiao Jin, Bhuwan Dhingra, Zhengping Liu, William Cohen, and Xinghua Lu. PubMedQA: A dataset for biomedical research question answering. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing*, pages 2567–2577, 2019.

Mandar Joshi, Eunsol Choi, Daniel S. Weld, and Luke Zettlemoyer. TriviaQA: A large scale distantly supervised challenge dataset for reading comprehension. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics*, pages 1601–1611, 2017.

Vladimir Karpukhin, Barlas Oğuz, Sewon Min, Patrick Lewis, Ledell Wu, Sergei Edunov, Danqi Chen, and Wen-tau Yih. Dense passage retrieval for open-domain question answering. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing*, pages 6769–6781, 2020.

Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris Alberti, Danielle Epstein, Illia Polosukhin, Jacob Devlin, Kenton Lee, Kristina Toutanova, Llion Jones, Matthew Kelcey, Ming-Wei Chang, Andrew M. Dai, Jakob Uszkoreit, Quoc Le, and Slav Petrov. Natural questions: A benchmark for question answering research. *Transactions of the Association for Computational Linguistics*, 7:453–466, 2019.

Philippe Laban, Tobias Schnabel, Paul N. Bennett, and Marti A. Hearst. SummaC: Re-visiting NLI-based models for inconsistency detection in summarization. In *Transactions of the Association for Computational Linguistics (TACL)*, 2022.

Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems*, volume 33, pages 9459–9474, 2020.

Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. G-Eval: NLG evaluation using GPT-4 with better human alignment. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 2511–2522, 2023. URL https://arxiv.org/abs/2303.16634.

Joshua Maynez, Shashi Narayan, Bernd Bohnet, and Ryan McDonald. On faithfulness and factuality in abstractive summarization. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 1906–1919, 2020.

Sewon Min, Kundan Krishna, Xinxi Lyu, Mike Lewis, Wen-tau Yih, Pang Wei Koh, Mohit Iyyer, Luke Zettlemoyer, and Hannaneh Hajishirzi. FActScore: Fine-grained atomic evaluation of factual precision in long form text generation. *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 11532–11545, 2023.

Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. SQuAD: 100,000+ questions for machine comprehension of text. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing*, pages 2383–2392, 2016.

Parth Sarthi, Salman Abdullah, Aditi Tuli, Shubh Khanna, Anna Goldie, and Christopher D. Manning. RAPTOR: Recursive abstractive processing for tree-organized retrieval. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2024. URL https://arxiv.org/abs/2401.18059.

Fangyuan Xu, Weijia Shi, and Eunsol Choi. RECOMP: Improving retrieval-augmented LMs with compression and selective augmentation. In *Proceedings of the International Conference on Learning Representations (ICLR)*, 2024. URL https://arxiv.org/abs/2310.04408.

Shi-Qi Yan, Jia-Chen Gu, Yun Zhu, and Zhen-Hua Ling. Corrective retrieval augmented generation. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 2564–2579, 2024.

Zhilin Yang, Peng Qi, Saizheng Zhang, Yoshua Bengio, William W. Cohen, Ruslan Salakhutdinov, and Christopher D. Manning. HotpotQA: A dataset for diverse, explainable multi-hop question answering. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 2369–2380, 2018.